

Contents

List of Symbols

1	Introduction	1
2	Preliminaries	3
2.1	Notation	3
2.2	Nonsmooth analysis and optimization	3
3	Bundle Methods	6
3.1	A basic bundle method	6
3.1.1	Derivation of the bundle method	6
3.1.2	Aggregate objects	9
3.2	13
3.2.1	Nonconvex Bundle Methods with Exact Information	13
3.2.2	Nonconvex bundle methods	13
3.2.3	Convex Bundle Methods with Inexact Information	15
3.3	How to deal with inexact information in bundle methods?	15
3.4	Proximal bundle method for nonconvex functions with inexact information	17
3.4.1	New subsubsection?	17
3.4.2	if convex function	21
4	How is inexact information dealt with?	22
5	Extension with Limited Memory approach	22
5.0.3	Thoughts about line search	22
5.1	Convergence	24
6	Variable Metric Bundle Methods	25
6.1	Subproblem Variable Metric	25

References

1 Introduction

There exists a sound and board theory of classical nonlinear optimization. However, this theory puts strong differentiability requirements on the given problem. Requirements that cannot always be fulfilled in practice. Examples for such practical application reach from problems in physics and mechanical engineering [2] over optimal control problems up to data analysis [1] and machine learning [19]. Other possible fields of applications are risk management and financial calculations [15, 20]. Additionally there exist so called stiff problems that are theoretically smooth but numerically nonsmooth due to rapid changes in the gradient [11].

There exists therefore a need for nonsmooth, that is not necessarily differentiable, optimization algorithms. A lot of the underlying theory and was developed in the 1970's, also driven by the "First World Conference on Nonsmooth Optimization" taking place in 1977 [14]. Now, there exists a well understood theoretical framework of nonsmooth analysis to create the basis for practical algorithms.

The most popular methods to tackle nonsmooth problems at the moment are bundle methods [22]. First developed only for convex functions [10] the method was soon extended to cope also with nonconvex objective functions [13].

Some time later these algorithms were again enhanced to deal with inexact information of the function value, the subgradient or both.

Some natural applications for these cases are derivative free optimization and stochastic simulations [22]. **Some more examples from different sources? Bilevel Problems?**

The basic idea of bundle methods is to model the original problem by a simpler function, often some sort of stabilized cutting plane model, that is minimized as a subproblem of the algorithm [7].

Adapt this part to what I finally really do:

In this thesis two different types of model functions will be examined that allow the use of inexact information in small to medium-scale problems as well as in large-scale problems. A limited memory approach is examined for the latter case.

what new? Combination of large-scale and inexact information - why needed
don't forget What - why - how

Adapt this part to what I finally really do:

This thesis is organized as follows:

introduction of the most important definitions and results of nonsmooth analysis. Then the introduction of a very basic bundle algorithm which is then generalized for nonconvex functions with nonsmooth optimization.

Throughout study of this algorithm including comparison to other approaches to tackle inexact information.

Introduction of variable metric (bundle) algorithm to tackle large-scale applications. "discussion" how far this is compatible with inexactness.

Numerical testing

discussion

First a proximal bundle method **Difference between different regularizations explained before...** large-scale optimization: a metric bundle method instead of a proximal bundle method -> limited memory approach

from PhD-thesis

-

2 Preliminaries

Theoretical Background, nonsmooth Analysis ???

Check if requirements on functions are stated and defined.

2.1 Notation

Throughout this thesis I consider the optimization Problem

$$\min_x f(x), \quad x \in X \subseteq \mathbb{R}^n \quad (1)$$

where f is a possibly nonsmooth function. Also write something about inexactness? specify X more precisely? Convex?

When it comes to nonsmooth objective functions the derivative based framework of non-linear optimization methods does not work any more. Therefore the most important definitions and results needed when working with nonsmooth functions are stated in this section.

Just definition, lemma, theorem or a bit explanation around it?

better just in Text without Definition, ...

See if requirements in definitions and theorems meet what is needed/provided later.

2.2 Nonsmooth analysis and optimization

A necessary assumption on the objective function f is that it is locally Lipschitz. This assumption assures the well-definedness of the following generalizations of derivatives.

Definition 2.1. [12] A function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is called *locally Lipschitz* if it is Lipschitz on each bounded subset $B \subseteq \mathbb{R}^n$

$$|f(y) - f(x)| \leq C\|y - x\| \quad \forall x, y \in B, \quad C > 0.$$

All convex functions are locally Lipschitz [8].

For convex functions one can define so called subgradients as a generalization of the usual derivative. They are defined using the directional derivative.

Definition 2.2. [8] The *directional derivative* of a convex function f at x in direction d is

$$f'(x, d) := \lim_{\lambda \downarrow 0} \frac{f(x + \lambda d) - f(x)}{\lambda}.$$

Definition 2.3. [8] Let f convex. The *subdifferential* $\partial f(x)$ of f at x is the nonempty compact convex set

$$\partial f(x) = \{g \in \mathbb{R}^n | f'(x, d) \geq \langle g, d \rangle \forall d \in \mathbb{R}^n\}.$$

The subdifferential is a convex set, that supports the graph of the function f from below. If f is differentiable at the point x , the subdifferential reduces to the gradient at that point [8].

This concept was generalized by Clarke for nonconvex functions. First a generalization of the directional derivative is given:

Definition 2.4. [2] Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ locally Lipschitz. The *generalized directional derivative* of f at x in direction d is given by

$$f^\circ(x, d) := \limsup_{\substack{y \rightarrow x \\ \lambda \downarrow 0}} \frac{f(y + \lambda d) - f(y)}{\lambda}.$$

This allows for the following definition.

Definition 2.5. [2] The *generalized gradient* of the locally Lipschitz function f at x is a nonempty convex compact set $\partial f(x)$ given by

$$\partial f(x) := \{g \in \mathbb{R}^n | f^\circ(x, d) \geq \langle g, d \rangle \forall d \in \mathbb{R}^n\}.$$

!!!other definition -> take definition from rockefeller/Hare directly from Paper!!!

If f is a convex function the generalized gradient coincides with the subdifferential ∂f of f [2].

Why epsilon Subdifferential?

implementable stopping criterion

dual form of bundle algorithms = same as stopping criterion?

Definition 2.6. [12] The function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is *semismooth* at $x \in \mathbb{R}^n$ if f is Lipschitz on a ball $\mathbb{B}_\varepsilon(x)$ around x and for each $d \in \mathbb{R}^n$ and for any sequences $\{t_k\} \subseteq \mathbb{R}_+$, $\{\theta_k\} \subseteq \mathbb{R}^n$ and $\{g_k\} \subseteq \mathbb{R}^n$ such that

$$\{t_k\} \downarrow 0, \quad \{\theta_k/t_k\} \rightarrow 0 \in \mathbb{R}^n \quad \text{and} \quad g_k \in \partial f(x + t_k d + \theta_k),$$

the sequence $\{\langle g_k, d \rangle\}$ has exactly one accumulation point.

check if I need ∞ -functions and if something changes then ϵ -subdifferential continuity properties of generalized gradient?

definitions from chapter inexact information

definition of inexactness for nonconvex kind of generalization of ϵ -subdifferential for non-convex case (Noll, inex, nonconv)

3 Bundle Methods

When bundle methods were first introduced in 1975 by Claude Lemaréchal and Philip Wolfe they were developed to minimize a convex (possibly nonsmooth) function f for which at least one subgradient at any point x can be computed [14].

To provide an easier understanding of the proximal bundle method in [22] and stress the most important ideas of how to deal with nonconvexity and inexactness first a basic bundle method is shown here.

link to chapter?

Bundle methods can be interpreted in two different ways: From the dual point of view one tries to approximate the ε -subdifferential to finally ensure first order optimality conditions. The primal point of view interprets the bundle method as a stabilized form of the cutting plane method where the objective function is modeled by tangent hyperplanes [6]. I focus here on the primal approach.

In the next two sections the function f is assumed to be convex.

notation, definitions

already done in previous preliminaries chapter?

3.1 A basic bundle method

This section gives a short summary of the derivations and results of chapter XV in [7] where a primal bundle method is derived as a stabilized version of the cutting plane method. If not otherwise indicated the results in this section are therefore taken from [7].

The optimization problem considered in this section is

$$\min_x f(x) \quad \text{s.t.} \quad x \in X \tag{2}$$

with the convex function f and the closed and convex set $X \subseteq \mathbb{R}^n$.

Define Problem again?? Incorporate “set-constraint” by writing $h(x) := f(x) + \mathbb{I}_X$. → later???

explanation

3.1.1 Derivation of the bundle method

The geometric idea of the cutting plane method is to build a piecewise linear model of the objective function f that can be minimized more easily than the original objective function.

This model is built from a *bundle* of information that is gathered in the previous iterations.

In the k 'th iteration, the bundle consists of the previous iterates x^j , the respective function values $f(x^j)$ and a subgradient at each point $g^j \in \partial f(x^j)$ for all indices j in the index set J_k . From each of these triples, one can construct a linear function

$$l_j(x) = f(x^j) + (g^j)^\top (x - x^j) \quad (3)$$

with $f(x^j) = l_j(x^j)$ and due to convexity $f(x) \geq l_j(x)$, $x \in X$.

One can now model the objective function f by the piecewise linear function

$$m_k(x) = \max_{j \in J_k} l_j(x) \quad (4)$$

and find a new iterate x^{k+1} by solving the subproblem

$$\min_x m_k(x) \quad \text{s.t.} \quad x \in X. \quad (5)$$

This subproblem should of course be easier to solve than the original task. A question that depends a lot on the structure of X . If $X = \mathbb{R}^n$ or a polyhedron, the problem can be solved easily. Still there are some major drawbacks to the idea. For example if $X = \mathbb{R}^n$ the solution of the subproblem in the first iteration is always $-\infty$.

In general one can say that the subproblem does not necessarily have to have a solution. To tackle this problem a penalty term is introduced to the subproblem:

$$\min \tilde{m}_k(x) = m_k(x) + \frac{1}{2t} \|x - x^k\|^2 \quad \text{s.t.} \quad x \in X \quad (6)$$

This new subproblem is strongly convex and has therefore always a unique solution.

how much explanation here? $\max_{j \in J_k} l_j(\hat{x}^k + d)$

Some nice sentences to explain the term a little bit more and to lead over to the next paragraph.

To understand the deeper motivation of this term see [7]. For this introduction it suffices to see that due to the regularization term the subproblem is now strongly convex and therefore always uniquely solvable.

The second major step towards the bundle algorithm is the introduction of a so called *stability center* or *serious point* \hat{x}^k . It is the iterate that yields the “best” approximation of the optimal point up to the k 'th iteration (not necessarily the best function value though).

The updating technique for \hat{x}^k is crucial for the convergence of the method: If the next iterate yields a decrease of f that is “big enough”, namely bigger than a fraction of the decrease suggested by the model function for this iterate, the stability center is moved to that iterate. If this is not the case, the stability center remains unchanged.

In practice this looks the following:

Define first the *nominal decrease* δ_k which is the decrease of the model for the new iterate

x^{k+1} compared to the function value at the current stability center \hat{x}^k .

$$\delta_k = f(\hat{x}^k) - \tilde{m}_k(x^{k+1}) + a_k \geq 0 \quad (7)$$

The nominal decrease is in fact stated a little differently for different versions of the bundle algorithm, this is why I added the constant $a_k \in \mathbb{R}$ here for generalization. In practice the difference between the decreases is not influencing the algorithm as δ_k is weighted by the constant $m \in (0, 1)$ for the descent test which compensates a_k .

If the actual decrease of the objective function is bigger than a fraction of the nominal decrease

$$f(\hat{x}^k) - f(x^{k+1}) \geq m\delta_k, \quad m \in (0, 1)$$

set the stability center to $\hat{x}^{k+1} = x^{k+1}$. This is called a *serious* or *descent step*.

If this is not the case a *null step* is executed and the serious iterate remains the same $\hat{x}^{k+1} = \hat{x}^k$.

The subproblem can be rewritten as a smooth optimization problem. For convenience rewrite the affine functions l_j with respect to the stability center \hat{x}^k .

citation for this???!!!

$$l_j(x) = f(x^j) + g^{j\top}(x - x^j) \quad (8)$$

$$= f(\hat{x}^k) + g^{j\top}(x - \hat{x}^k) - (f(\hat{x}^k) - f(x^j) + g^{j\top}(x^j - \hat{x}^k)) \quad (9)$$

$$= f(\hat{x}^k) + g^{j\top}(x - \hat{x}^k) - e_j^k \quad (10)$$

where

$$e_j^k = f(\hat{x}^k) - f(x^j) + g^{j\top}(x^j - \hat{x}^k) \geq 0 \quad \forall j \in J_k \quad (11)$$

is the *linearization error*. The nonnegativity property is essential for the convergence theory and will also be of interest when moving on to the case of nonconvex and inexact objective functions.

Subproblem (6) can now be written as

$$\min_{\hat{x}^k + d \in X} \tilde{m}_k(d) = f(\hat{x}^k) + \max_{j \in J_k} \{g^{j\top}d - e_j^k\} + \frac{1}{2t_k} \|d\|^2 \quad (12)$$

$$\Leftrightarrow \min_{\hat{x}^k + d \in X, \xi \in \mathbb{R}} \xi + \frac{1}{2t_k} \|d\|^2 \quad \text{s.t.} \quad f(\hat{x}^k) + g^{j\top}d - e_j^k - \xi \leq 0, \quad j \in J_k \quad (13)$$

where the constant term $f(\hat{x}^k)$ was discarded for the sake of simplicity.

If X is a polyhedron this is a quadratic optimization problem that can be solved using

standard methods of nonlinear optimization. The pair (ξ_k, d^k) solves (13) if and only if d^k solves the original subproblem (12) and $\xi_k = f(\hat{x}^k) + \max_{j \in J_k} g^j{}^\top d^k - e_j^k$. The new iterate is then given by $x^{k+1} = \hat{x}^k + d^k$. **citation!!!**

Remark: Setting $\check{f}(x) = f(x) + \mathbb{I}_X(x)$ the above optimization problem is ...
The proximal point mapping or prox-operator

$$\text{prox}_{t,f}(x) = \arg \min_y \left\{ \check{f}(y) + \frac{1}{2t} \|x - y\|^2 \right\}, \quad t > 0 \quad (14)$$

source??? This special form of the subproblems gives the proximal bundle method its name and will occur again later???

3.1.2 Aggregate objects

The constraint $\hat{x}^k + d \in X$ can also be incorporated directly in the objective function by using the indicator function

$$\mathbb{I}_X(x) = \begin{cases} 0, & \text{if } x \in X \\ +\infty, & \text{if } x \notin X \end{cases}.$$

Subproblem (6) then writes as

$$\min_{\hat{x}^k + d \in \mathbb{R}^n, \xi \in \mathbb{R}} \xi + \mathbb{I}_X + \frac{1}{2t_k} \|d\|^2 \quad \text{s.t.} \quad g^j{}^\top d - e_j^k - \xi \leq 0, \quad j \in J_k \quad (15)$$

check if f also not put into subproblem before

Some introduction how this and the aggregate error expression relate to each other. Why it is in this case easier to write the model in the nonsmooth form...

Lemma XI 3.1.1 $\partial g = \partial f + \partial \mathbb{I}_X$ for $g = f + \mathbb{I}_X$.

One gets the following results about the step d^k of the subproblem:

Lemma 3.1. The optimization problem (15) has for $t_k > 0$ a unique solution given by

$$d^k = -t_k(G^k + \nu^k), \quad G^k \in \partial m_k(d^k), \quad \nu^k \in \partial \mathbb{I}_X. \quad (16)$$

Furthermore

$$m_k(\hat{x}^k + d) \geq f(\hat{x}^k) + G^k{}^\top d - E_k \quad \forall d \in \mathbb{R}^n \quad (17)$$

inequality because of aggregation technique. Is sharp when cutting plane model is used?
source?

where

$$E_k := f(\hat{x}^k) - m_k(x^{k+1}) + G^k{}^\top d^k. \quad (18)$$

Comment on the inequality missing

The quantities G^k and E^k are the *aggregate subgradient* and the *aggregate error*.

Explain aggregation process in more detail

From the Karush-Kuhn-Tucker conditions (KKT-conditions) one can see that in the optimum there exist Lagrange or *simplicial multiplier* α_j^k , $j \in J_k$ such that

$$\alpha_j^k \geq 0, \quad \sum_{j \in J_k} \alpha_j^k = 1 \quad (19)$$

by rewriting and so on... one can see that the above expressions are in fact

From the dual problem one obtains that the aggregate subgradient and error can also be expressed as

$$E_k = \sum_{j \in J_k} \alpha_j^k e_j^k \quad \text{and} \quad G^k = \sum_{j \in J_k} \alpha_j^k g^j. \quad (20)$$

Finally use Lemma ??? in [7]

$$m_k(x^{k+1}) = f(\hat{x}^k) - E_k - t_k \|G^k\|^2$$

to reformulate the nominal decrease δ_k :

$$\delta_k = f(\hat{x}^k) - m_k(x^{k+1}) - \frac{1}{2} t_k \|G^k\|^2 = E_k + \frac{1}{2} t_k \|G^k\|^2$$

The nominal decrease in this case is defined as:

noch mal anschauen

$$\delta_k := E_k + t_k \|G^k + \nu^k\|^2 = f(\hat{x}^k) - m_k(x^{k+1}) - \nu^{k\top} d^k \quad (21)$$

In practice the different definition of the decreases makes no difference because of the weighting with the descent parameter m .

The following basic bundle algorithm can now be stated:

Reformulate equations, model function

introduce aggregate expressions

say something to J -update, say something to t -update

see if all abbreviations (f_j, g^j, \dots) are introduced

introduce prox-operator and proximal points

algorithm

Basic bundle method

Select descent parameter $m \in (0, 1)$ and a stopping tolerance $\text{tol} \geq 0$. Choose a starting point $x^1 \in \mathbb{R}^n$ and compute $f(x^1)$ and g^1 . Set the initial index set $J_1 := \{1\}$ and the initial stability center to $\hat{x}^1 := x^1$, $f(\hat{x}^1) = f(x^1)$ and select $t_1 > 0$.

For $k = 1, 2, 3 \dots$

1. Calculate

$$d^k = \arg \min_{d \in \mathbb{R}^n} m_k(\hat{x}^k + d) + \mathbb{I}_X + \frac{1}{2t_k} \|d\|^2$$

and the corresponding Lagrange multiplier α_j^k , $j \in J_k$. say how model m_k looks here. include \mathbb{I}_X

2. Set

$$G^k = \sum_{j \in J_k} \alpha_j^k g_j^k, \quad E_k = \sum_{j \in J_k} \alpha_j^k e_j^k, \quad \text{and} \quad \delta_k = E_k + t_k \|G^k + \nu^k\|^2$$

If $\delta_k \leq \text{tol} \rightarrow \text{STOP}$.

3. Set $x^{k+1} = \hat{x}^k + d^k$.

4. Compute $f(x^{k+1})$, g^{k+1} .

If

$$f^{k+1} \leq \hat{f}^k - m\delta_k \rightarrow \text{serious step.}$$

Set $\hat{x}^{k+1} = x^{k+1}$, $f(\hat{x}^{k+1}) = f(x^{k+1})$ and select suitable $t_{k+1} > 0$.

Otherwise \rightarrow nullstep.

Set $\hat{x}^{k+1} = \hat{x}^k$, $f(\hat{x}^{k+1}) = f(x^{k+1})$ and choose t_{k+1} in a suitable way.

5. Select new bundle index set $J_{k+1} = \{j \in J_k | \alpha_j^{k+1} \neq 0\} \cup k+1$, calculate e_j for $j \in J_{k+1}$ and update the model m_k .
-

In steps 4 and 5 of the algorithm the updates of the steplength t_k and the index set J_k are only given in a very general form.

The “suitable” choice of t_k will be discussed more closely in the convergence analysis of decide which method; say that $t_k > 0 \forall k \dots$

Comment on J_k update \rightarrow depends on what is included in thesis.

For the choice of the new index set J_{k+1} different aggregation methods to keep the memory size controllable are available. The most easy and intuitive one is to just take those parts of the model function, that are actually active in the current iteration. This is done in this basic version of the method.

Refer to low memory bundling if later in thesis. Instead of keeping every index in the set J_k different compression ideas exist. For now I therefor stick to this update.

refer to later “low memory” thing??

explanation to t_k update. \rightarrow include at which point??? This simple idea has however some major drawbacks [8]:

- Minimization of the cutting plane model of the objective function is not trivial. Indeed unconstrained minimization of the model is never possible in the first step, where it is just a line, unless the starting point is already a minimum.
- The convergence speed is very slow.

If convergence speed named here, does it have to be shown (rates)? For all algorithms???

Leave out? Argue about instability?

To address those issues a regularization is added to the cutting plane model. This ensures unique solvability of the minimization of the subproblem. By introducing a stability center and

3.2 ...

possible simplifications of the algorithm

3.2.1 Nonconvex Bundle Methods with Exact Information

Simplification / better results if exact information The main ideas of the algorithm are basically the ones developed in [6] for the redistributed proximal bundle method for exact nonconvex problems.

Setting the error bounds $\bar{\sigma}$ and $\bar{\theta}$ to zero results therefore in the following convergence theorem.

Theorem 3.2. *Let the sequence $\{\eta_k\}$ be bounded, $\liminf_{k \rightarrow \infty}$ and the cardinality of the set $\{j \in J_k | \alpha_j^k > 0\}$ be uniformly bounded in k .*

Then every accumulation point of sequence of serious iterates $\{\hat{x}^k\}$ is a stationary point of the problem.

think last condition only interesting in inexact case.

try to gain some insight with generalized ε -subdifferential from Chinese paper:

ε -limiting subdifferential []

In the exact case boundedness of the sequence $\{\eta_k\}$ is proven for lower- \mathcal{C}^2 functions in [6]. This is not possible in the inexact case, even if the objective function f is convex.

A further simplification of the method for exact information is not necessary as the method is already almost as simple as the basic bundle method for nonconvex exact functions. Additionally no new concepts needed to be introduced when doing the step from nonconvex exact problems, for which the algorithm was originally designed, to problems with inexact information.

Remark: I want to add here, that the simplicity of the algorithm is rather special for methods suitable for nonconvex problems. Often a linesearch algorithm has to be inserted in the nonconvex case, which is not needed here.

3.2.2 Nonconvex bundle methods

There are different approaches for handling nonconvexity of the objective function in bundle methods. As the nonnegativity property of the linearization errors e_j^k is crucial for the convergence proof of convex bundle methods an early idea was forcing the errors to be so by different downshifting strategies. A very common one is using the *subgradient locality measure* [9, 13]. Here the linearization error is essentially replaced by the nonnegative number

$$\tilde{e}_j^k := \max_{j \in J_k} \{|e_j^k|, \gamma \|\hat{x}^k - x^j\|^2\} \quad (22)$$

or a variation of this expression.

Remark on dual view? How subgradient locality measure measures how close subgradient is to subdifferential of f ???

Methods using this kind of manipulation of the model function are often endowed with a line search to provide sufficient decrease of the objective function. For the linesearch to terminate finitely, semismoothness of the objective function is usually needed.

It can be proven that every accumulation point of the sequence of serious points $\{\hat{x}^k\}$ is a stationary point of the objective function f under the additional assumptions that f is locally Lipschitz and the level set $\{x \in \mathbb{R}^n | f(x) \leq f(\hat{x}^1)\}$ is bounded [5].

A drawback to the method described above is that it is primarily supported from the dual point of view of the bundle algorithm. Newer concepts focus also on the primal point of view. This invokes for example having different model functions for the subproblem.

In [3, 4] the difference function

$$h(d) := f(x^j + d) - f(x^j) \quad j \in J_k \quad (23)$$

is approximated to find descent direction of f .

The negative linearization errors are addressed by having two different bundles. One containing the indices with nonnegative linearization errors and one containing the other ones. From these two bundles two cutting plane approximations can be constructed which provide the bases for the calculation of the new iterate.

Convergence of the method to a stationary point is proven under the assumption of f being locally Lipschitz and semismooth.

still line search needed

any bounded (level-)sets needed???

In [17] Noll et al. follow an approach of approximating a local model of the objective function. The model can be seen as a nonsmooth generalization of the Taylor expansion and looks the following:

$$\Phi(y, x) = \phi(y, x) + \frac{1}{2}(y - x)^\top Q(x)(y - x) \quad (24)$$

The so called *first order model* $\phi(\cdot, x)$ is convex but possibly nonsmooth and can be approximated by cutting planes. The *second order part* is a quadratic but not necessarily convex. The algorithm then proceeds a lot in the lines of a general bundle algorithm.

The method relies on a smart management of the proximity parameter τ_k which corresponds to $1/t_k$ in the notation of this thesis. This is why the method does not need a linesearch subroutine. For a locally Lipschitz objective function with a bounded levelset $\{x \in \mathbb{R}^n | f(x) \leq f(\hat{x}^1)\}$ convergence to a stationary point is established.

In paper [17] stated that proximity control = proximal bundle Algorithm with smart

t_k -control very powerful

Add Luksan in view of Karmita Method? Then short introduction of variable metric bundle algorithms necessary; would be manageable in this section

For proximal bundle methods: two strategies: line search or (newer) proximity control: It seems that a successful strategy to deal with nonconvexity is proximity control as used in different manners in [?, ?, ?, ?, 17, ?, ?]

3.2.3 Convex Bundle Methods with Inexact Information

- stronger convergence results possible because of exploitation of convexity
- changes in the algorithm because if convexity should be exploited: inexactness cannot be treated as nonconvexity
-

in extra section??

3.3 How to deal with inexact information in bundle methods?

Partition section in two parts:

1. How is generally dealt with inexact information in the algorithms
- 2 a) What information is inexact (only subgradients/both...) → what do you gain from this?
- 2 b) What kind of assumptions are on the inexactness? (asymptotic, only over- /under-estimation?)

- recognized: fundamentally different? approach for convex and nonconvex functions (at least in algorithm)
convex: “deal” with inexactness; extra steps...
nonconvex: generally no difference in algorithm (but for example line search not possible → only no change, if algorithm was suitable before)
- nonconvex algorithms: inexactness is seen as some kind of nonconvexity → for function values clear, for subgradients???
- in convex case: often assumption, that gradient is from ε -subdifferential
is this restrictive? →

What does “approximate subgradient” mean???

generally seems to be that for convex functions it is the same concept whether one takes the ε -subdifferential or a ball around the regular subdifferential.

seems to be the same:

$$\|g_a - g\| \leq \theta \quad (25)$$

$$\Leftrightarrow g_a \in \partial f + B_\theta(0) \quad (26)$$

$$\Leftrightarrow g_a \in \partial_\varepsilon f, \quad \theta \leq \varepsilon^2 \quad (27)$$

Last implication only for convex functions because ε -subdifferential otherwise not defined. See also papers from “Chinese-search”

Different “degrees” of inexactness: inexact subgradients; also function values (only subgradients easier??); asymptotically exact; exactness only for serious steps, not at null steps; accuracy controllable or not \rightarrow throughout study in in depth paper.

One can clearly see, that at the moment there exist two fundamentally different approaches to tackle inexactness in various bundle methods depending on if the method is developed for convex or nonconvex objective functions.

In the nonconvex case inexactness is only considered in the paper by Hare, Sagastizàbal and Sodolov [22] presented above and Noll [16]. In these cases the inexactness can be seen as an “additional nonconvexity”. In practice this means that the algorithm can be taken from the nonconvex case with no or only minor changes.

In case of convex objective functions changes in the algorithm are more involved. The reason for this is that generally stronger convergence results are possible with inexactness in the convex case than in the nonconvex case. This means however, that the inexactness cannot be incorporated as easily into the algorithm.

Remark on nonconvexity line search and inexactness

A possible reason why there are not already more publication on bundle methods with inexact information in the nonconvex case although there exists a broad variety of algorithms that deal with the exact case could be that many of them incorporate a line search. To make sure that this subalgorithm is finite the objective function has to be semi-smooth **definition of semismoothness, check** This however cannot be the case when the functions values of the objective function are only approximated.

3.4 Proximal bundle method for nonconvex functions with inexact information

introduction

This section focuses on the proximal bundle method presented in [22].

The idea is to extend the basic bundle algorithm for nonconvex functions with both inexact function and subgradient information.

The key idea of the algorithm is the one already developed for [6]: When dealing with nonconvex functions a very critical difference to the convex case is that the linearization errors are not necessarily nonnegative any more. To tackle this problem the errors are manipulated to enforce nonnegativity. In this case this is done by modeling not the objective function directly but a convexified version of it.

3.4.1 New subsection?

“assumptions and notations”

introduce exact optimization problem that is used in this section and its properties if not already introduce in “Preliminaries”.

Throughout this section the optimization problem

$$\min_x f(x) \quad \text{s.t.} \quad x \in X \quad (28)$$

where f is locally Lipschitz is considered. $X \subseteq \mathbb{R}^n$ is assumed to be a convex compact set. Both the function value as well as the subgradient can be provided in an inexact form.

For the function value inexactness is defined straight forwardly: If

$$\|\tilde{f} - f(x)\| \leq \sigma \quad (29)$$

then \tilde{f} approximates the value $f(x)$ within σ .

For the subgradients inexactness is interpreted in the following way: $\tilde{g} \in \mathbb{R}^n$ approximates a subgradient $g \in \partial f(x)$ within $\theta \geq 0$ if

$$\tilde{g} \in \partial f(x) + B_\theta(0). \quad (30)$$

In the paper it is assumed that the errors are bounded although the bound does not have to be known.

$$|\sigma_j| \leq \bar{\sigma} \quad \text{and} \quad 0 \leq \theta_j \leq \bar{\theta} \quad \forall j \in J_k. \quad (31)$$

In the context of inexact information it is important to make a distinction between

the (unknown) exact function value and its approximation. Throughout this chapter I therefore write $f(x)$ for the exact function value whereas the approximation will be written as f_j or \hat{f}_k for the approximation at the current stability center.

The objective function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is assumed to be proper, (subdifferentially) regular and locally Lipschitz continuous with full domain.

Definition 3.3. [18] A function $f : \mathbb{R}^n \rightarrow \bar{\mathbb{R}} = [-\infty, +\infty]$ is called *proper* if $f(x) < \infty$ for at least one $x \in \mathbb{R}^n$ and $f(x) > \infty \forall x \in \mathbb{R}^n$.

Definition 3.4. [18] $f : \mathbb{R}^n \rightarrow \bar{\mathbb{R}}$ is called *subdifferentially regular* at \bar{x} if $f(\bar{x})$ is finite and the epigraph

$$\text{epi}(f) := \{(x, \alpha) \in \mathbb{R}^n \times \mathbb{R} \mid \alpha \geq f(x)\}$$

is Clarke regular at $\bar{x}, f(\bar{x})$.

Closed convex sets are Clarke regular, so in particular the epigraph of lower \mathcal{C}^2 -functions?.

Definition semismooth for later:

Definition 3.5. A function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is called *semismooth* at $x \in \mathbb{R}^n$ if f is Lipschitz near x and for each $d \in \mathbb{R}^n$ and for any sequences $\{t_k\} \subseteq \mathbb{R}_+, \{\theta^k\} \subseteq \mathbb{R}^n$ and $\{g^k\} \subseteq \mathbb{R}^n$ such that

$$\{t_k\} \downarrow 0, \quad \{\theta^k/t_k\} \rightarrow 0 \in \mathbb{R}^n \quad \text{and} \quad g^k \in \partial f(x + t_k d + \theta^k),$$

the sequence $\{\langle g^k, d \rangle\}$ has exactly one accumulation point.

Definition 3.6. A point $x \in \mathbb{R}^n$ that satisfies $0 \in \partial f(x)$ is called a *stationary point* of f .

explanation

A main issue both nonconvexity and inexactness entail is that the linearization errors e_j^k are not necessarily nonnegative any more.

So based on the results in [21] not the objective function but a convexified version of it is modeled as the objective function of the subproblem.

When looking at the subproblem formulated as in (12) one can see that the new iterate x^{k+1} is in fact a *proximal point* of the subproblem.

The *proximal point mapping* or *prox-operator* is defined as

$$\text{prox}_{t,f}(x) = \arg \min_y \left\{ \check{f}(y) + \frac{1}{2t} \|x - y\|^2 \right\}, \quad t > 0 \quad (32)$$

For $\check{f}(x) := m(x) + \mathbb{I}_X(x)$ and $\mu := \frac{1}{t_k}$ this is just subproblem (12) with the constraint $x \in X$ incorporated in the objective function. Because of this special form of the subproblems primal bundle methods are also called proximal bundle methods.

explain in much more detail when read about calculation of proximal points for nonconvex functions. At the moment just main ideas.

The key idea is now to use the relation

$$\text{prox}_{R=\mu+\eta, f}(x) = \text{prox}_{\mu, f+\eta/2 \cdot \|\cdot-x\|^2}(x). \quad (33)$$

This means, that the proximal point of the function f for parameter $R = \eta + \mu$ is the same as calculating the proximal point of the regularized function

$$\tilde{f}(y) = f(y) + \frac{\eta}{2} \|y - x\|^2 \quad (34)$$

with respect to the parameter μ . η is therefore called the *convexification parameter* and μ is the *prox-parameter*.

So the function that will be modeled by the cutting plane approximation is no longer the original objective function f but the convexified version \tilde{f} .

The linear functions forming the model have therefore a tilted slope

$$s_j^k = g^j + \eta_k (x^j - \hat{x}^k). \quad (35)$$

η is defined to be such that the augmented linearization error is nonnegative:

$$\eta_k \geq \max \left\{ \max_{j \in J_k, x^j \neq \hat{x}^k} \frac{-2e_j^k}{\|x^j - \hat{x}^k\|^2}, 0 \right\} + \gamma \quad (36)$$

With the “saveguarding parameter” $\gamma \geq 0$

explain, why linearization errors are defined like that

$$0 \leq c_j^k := e_j^k + b_j^k, \quad \text{with} \quad \begin{cases} e_j^k := \hat{f}_k - f_j - \langle g^j, \hat{x}^k - x^j \rangle \\ b_j^k := \frac{\eta_k}{2} \|x^j - \hat{x}^k\|^2 \end{cases} \quad (37)$$

The new model function can therefore be written as

$$M_k(d) := \hat{f}_k + \max_{j \in J_k} \{s_j^{k^\top} d - c_j^k\} \quad (38)$$

The definition of the aggregate objects follows straightforward:

$$S^k := \sum_{j \in J_k} \alpha_j^k s_j^k \quad (39)$$

$$C_k := \sum_{j \in J_k} \alpha_j^k c_j^k \quad (40)$$

explain how δ_k is derived.

$$\delta^k := C_k + t_k \|S^k + \nu^k\|^2 \quad (41)$$

algorithm

Nonconvex proximal bundle method with inexact information

Select parameters $m \in (0, 1)$, $\gamma > 0$ and a stopping tolerance $\text{tol} \geq 0$.

Choose a starting point $x^1 \in \mathbb{R}^n$ and compute f_1 and g^1 . Set the initial index set $J_1 := \{1\}$ and the initial prox-center to $\hat{x}^1 := x^1$, $\hat{f}_1 = f_1$ and select $t_1 > 0$.

For $k = 1, 2, 3, \dots$

1. Calculate

$$d^k = \arg \min_{d \in \mathbb{R}^n} \left\{ M_k(\hat{x}^k + d) + \mathbb{I}_X(\hat{x}^k + d) + \frac{1}{2t_k} \|d\|^2 \right\}.$$

2. Set

$$\begin{aligned} G^k &= \sum_{j \in J_k} \alpha_j^k s_j^k, \quad \nu^k = -\frac{1}{t_k} d^k - G^k \\ C_k &= \sum_{j \in J_k} \alpha_j^k c_j^k \\ \delta_k &= C_k + t_k \|G^k + \nu^k\|^2 \end{aligned}$$

If $\delta_k \leq \text{tol} \rightarrow \text{STOP}$.

3. Set $x^{k+1} = \hat{x}^k + d^k$.

4. Compute f^{k+1}, g^{k+1}

If

$$f^{k+1} \leq \hat{f}^k - m\delta_k \rightarrow \text{serious step}$$

Set $\hat{x}^{k+1} = x^{k+1}$, $\hat{f}^{k+1} = f^{k+1}$ and select $t_{k+1} > 0$.

Otherwise \rightarrow nullstep

Set $\hat{x}^{k+1} = \hat{x}^k$, $\hat{f}^{k+1} = f^{k+1}$ and choose $0 < t_{k+1} \leq t_k$.

5. Select new bundle index set J_{k+1} , keeping all active elements. Calculate

$$\eta_k \geq \max \left\{ \max_{j \in J_{k+1}, x^j \neq \hat{x}^{k+1}} \frac{-2e_j^k}{|x^j - \hat{x}^{k+1}|^2}, 0 \right\} + \gamma$$

and update the model M^k

3.4.2 if convex function

Convergence for inexact convex functions:

- states in paper [22] (p. 14) that for convex functions error of $\bar{\sigma}$ instead of $2\bar{\sigma}$ possible (and for lower models; see depth paper?)

To Do:

- proof serious steps
- proof null steps
- limit of G^k
- proof in book; see if possible to leave out D ; compare should be possible if bounded level sets assumed; check this! compare with “depth” $\rightarrow \phi$
- see if η_k can be bounded in exact case - yes for the class of functions mentioned in the paper
- find counterexample, that η can't be bounded in inexact case??? - main argument: have to assure, that convexified objective function is “convex on all bundle points x^j ” from a certain η^k on
- compare nonconvex exact \leftrightarrow inexact convergence results only look at exact paper again if better results!
 - check if correct: inexact more general because choice of t_k more freely??? in exact μ_k only changed when restart update strategy not important for convergence; maybe for convergence speed?
 - check if update strategy important for convergence speed? - yes see napsu ... “Comparison ...”
- check if (ii) in Theorem 6 in paper can be assured by choice of t_k in algorithm (think yes)
- compare to convergence results of other papers check if better results can be carried over - results all the same; check for prerequisites on functions
- check if other papers have better prerequisites check if results can be carried over all need locally Lipschitz and either a compact? subset or bounded lower level sets Better in other papers?: neither $\{j \in J_k | \alpha_j^k > 0\}$ nor $\{\eta_k\}$ need to be bounded?? any prerequisites on t_k in other papers???
- most other solutions for nonconvex functions: based on dual idea; remark on dual

view on bundle method?

- write that down nicely
- read depth paper
 - see what kind my algorithm is
 - check Lemma 5 (iv) stronger?
- remark on that inexactness makes objective nonconvex \rightarrow relate to nonconvexity paper?
- check again uniformly bounded J_k
aggregated objects (3.4), 5.2, 7.1 in “depth”
- generalized gradients may only be shifted not tilted?
is this realistic assumption? \rightarrow all ok, tilted and shifted!
- ask Simon to ε -subdifferentials
- Algorithm:
 - print something useful in every iteration
 - η should not get too big

If (newer) papers needed: look at citations of the ones I have.

(sub-)Level sets of continuous functions are closed (image of closed set closed under continuous function); should also hold for lower semicontinuous functions in metrizable(???) space

4 How is inexact information dealt with?

5 Extension with second Order Models

Variable metric methods are also known as quasi-Newton methods.

PhD Thesis (p. 33) inheritance of positive definiteness of matrix update formulas

p. 34: number of correction pairs usually $3 \leq m \leq 30$

5.0.3 Thoughts about line search

- after what is written in “nonconv, inex”-paper: Line search not provable to be finite if inexact information
- is line search standard in all variable metric methods?
 - looks like it

- there do exist versions without line search \rightarrow other update???
- does it work without line search??? -think yes
- does prox-parameter t_k have some relation to linesearch/stepsize?

Algorithm 1 in [5]

Variable metric limited memory bundle method with inexact information

Select parameters $m \in (0, 1)$, $\gamma > 0$, $K > 0$, $\rho \in (0, \frac{1}{2})$ and a stopping tolerance $\text{tol} \geq 0$. Set the initial metric matrix $D_1 = \mathbb{I}_{n \times n}$. Choose a starting point $x^1 \in \mathbb{R}^n$ and compute f_1 and $s^1 = g^1 = S^1$. Set the initial serious iterate $\hat{x}^1 := x^1$, $\hat{f}_1 = f_1$, and $\hat{s}^1 = s^1$, and $c_1 = C_1 = 0$.

Set the correction indicator and the correction indicator for consecutive null steps $i_C = 0$.

For $k = 1, 2, 3, \dots$

1. Compute

$$d^k = -D_k S^k$$

by using a limited memory BFGS update if the step before was a serious step and by using a limited memory SR1 update otherwise.

For $k = 1$: $d^1 = -S^1$.

2. If $-S^{k\top} d^k < \rho S^{k\top} S^k$ or $i_{CN} = 1$ then set

$$d^k = d^k - \rho S^k, \quad (42)$$

and $i_C = 1$. If the previous step was a null step set also $i_{CN} = 1$.

Otherwise set $i_C = 0$.

3. Set

$$\delta_k^w = -S^{k\top} d^k + 2C_k \quad \text{and} \quad (43)$$

$$\delta_k^q = \frac{1}{2} S^{k\top} S^k + C_k. \quad (44)$$

If $\delta_k^w < \text{tol}$ and $\delta_k^q < \text{tol}$ stop with \hat{x}_k as the final solution.

4. Set

$$\theta_k = \min\{1, K/\|d^k\|\} \quad (45)$$

$$x^{k+1} = \hat{x}^k + \theta_k d^k. \quad (46)$$

Calculate the inexact function value f_{k+1} and an inexact subgradient g^{k+1} .

5. If

$$f^{k+1} \leq \hat{f}^k - m \delta_k^w \rightarrow \text{serious step}$$

Which δ ???

Set $\hat{x}^{k+1} = x^{k+1}$, $\hat{f}_{k+1} = f_{k+1}$, $s^{k+1} = S^{k+1} = g^{k+1}$.

Otherwise \rightarrow nullstep

Compute the convexification parameter

$$\eta_{k+1} = \max\{e_{k+1}, 0\} + \gamma$$

with $e_{k+1} = \hat{f}_{k+1} - f_{k+1} + \langle g^{k+1}, d^k \rangle$.

Set

$$\hat{x}^{k+1} = \hat{x}^k \tag{47}$$

$$\hat{f}^{k+1} = \hat{f}^k \tag{48}$$

$$s^{k+1} = g^{k+1} + \eta_{k+1} * (x^{k+1} - \hat{x}^{k+1}) \tag{49}$$

$$c_{k+1} = \max\{e_{k+1}, 0\} + \frac{\gamma}{2}. \tag{50}$$

Compute the new correction pair $u_1^k = \theta_k d^k$ and $u_2^k = s^{k+1} - \hat{s}^{k+1}$.

6. If this step was a serious step, go to 1.

In case of a null step determine multipliers $\alpha_i^k \geq 0$, $i = \{1, 2, 3\}$, $\sum_i \alpha_i^k = 1$ that minimize the function

$$\begin{aligned} \phi(\alpha_1, \alpha_2, \alpha_3) = & (\alpha_1 \hat{s}^{k+1} + \alpha_2 s^{k+1} + \alpha_3 S^k) D_k (\alpha_1 \hat{s}^{k+1} + \alpha_2 s^{k+1} + \alpha_3 S^k) \\ & + 2(\alpha_2 c_{k+1} + \alpha_3 C_k) \end{aligned}$$

where D_k is calculated by the same updating formula as in step 1 and $D_k = D_k + \rho \mathbb{I}$ if $i_C = 1$.

Compute the aggregate subgradient and error

$$S^{k+1} = \alpha_1 \hat{s}^{k+1} + \alpha_2 s^{k+1} + \alpha_3 S^k \tag{51}$$

$$C_{k+1} = (\alpha_2 c_{k+1} + \alpha_3 C_k) \tag{52}$$

and go back to step 1.

5.1 Convergence

6 ???

6.1 Subproblem Variable Metric

For comparison: Subproblem proximal bundle

$$\min_{d \in \mathbb{R}^n, \xi \in \mathbb{R}} \xi + \frac{1}{2t_k} \|d\|^2 = \xi + \frac{1}{2} d^\top \left(\frac{1}{t_k} \mathbf{I} \right) d \quad (53)$$

$$\text{s.t.} \quad f(\hat{x}^k) + g^j{}^\top d - e_j^k - \xi \leq 0, \quad j \in J_k \quad (54)$$

Subproblem variable metric:

$$\min_{d \in \mathbb{R}^n, \xi \in \mathbb{R}} \xi + \frac{1}{2} d^\top D_k d \quad (55)$$

$$\text{s.t.} \quad f(\hat{x}^k) + g^j{}^\top d - e_j^k - \xi \leq 0, \quad j \in J_k \quad (56)$$

These are \mathbb{R}^{n+1} dimensional quadratic optimization problems.

Find out if D_k is diagonal matrix! Think not.

Approaches not so different. Instead of just scaling the identity \rightarrow induce “curvature information” via past subgradients.

Dual proximal subproblem:

$$\min_{\alpha \in \mathbb{R}^{|J_k|}} \frac{1}{2} \left(\sum_{j \in J_k} \alpha_j g^j \right)^\top t_k \mathbf{I} \left(\sum_{j \in J_k} \alpha_j g^j \right) + \sum_{j \in J_k} \alpha_j e_j^k \quad (57)$$

$$\text{s.t.} \quad \sum_{j \in J_k} \alpha_j = 1 \text{ and } \alpha_j \geq 0 \quad j \in J_k \quad (58)$$

Dual variable metric subproblem:

$$\min_{\alpha \in \mathbb{R}^{|J_k|}} \frac{1}{2} \left(\sum_{j \in J_k} \alpha_j g^j \right)^\top D_k^{-1} \left(\sum_{j \in J_k} \alpha_j g^j \right) + \sum_{j \in J_k} \alpha_j e_j^k \quad (59)$$

$$\text{s.t.} \quad \sum_{j \in J_k} \alpha_j = 1 \text{ and } \alpha_j \geq 0 \quad j \in J_k \quad (60)$$

These are $\mathbb{R}^{|J_k|}$ dimensional quadratic optimization problems.

check linear independent g^j 's.

7 Application to Model Selection for Primal SVM

7.1 Introduction

A natural application for the inexact bundle algorithm is an optimization problem where the objective function is again an the result of an optimization. If this objective function is iteratively optimized, the function value is never exact. But by the stopping condition it is (often???) possible to have the needed error bound on the function value (what about the gradient???). (Comment on asymptotical exactness?)

An example for such a setting is bilevel optimization. A general bilevel problem can be formulated as (citation)

bilevel problem.

It consists of an *upper* or *outer level* which is the overall function to be optimized. Contrary to usual constrained optimization problems which are constrained by explicitly given equalities and inequalities a bilevel program is additionally constrained to a second optimization problem, the *lower* or *inner level* problem.

The solution of bilevel problems can be divided roughly in two classes: implicit and explicit solution methods. In the explicit methods the lower level problem is usually rewritten by its KKT conditions and the upper and lower level are solved simultaneously. In the setting described here this method is described in detail in [?]. The second approach is the implicit one. Here the lower level problem is solved in every iteration and the solution is plugged into the upper level objective. Obviously if the inner level problem is solved numerically, the solution cannot be exact. Additionally the **comment on solution map and nonsmoothness -> citation** This is why the inexact bundle algorithms seems a natural choice to tackle this problem. Comment on Moore-Paper??? write a little bit more explicit what kind of problem is tackled here (Model selection for support vector machines (machine learning) in a bilevel formulation)?????

7.2 Explanation SVM

Some introductory words on machine learning in general? Support vector machines...

7.3 Explanation Bilevel Approach and Inexact Bundle Method

Check if regression case also included????? The parameter in the objective function of the classification problem has to be set before hand. This step is part of the model selection process (citation) goal: set this parameter optimally A very intuitive and widely used approach: grid search (description) -> very costly, discrete parameter choice, not practicable in case of many parameter A more recent approach is the formulation as a bilevel problem used in [?, ?].

7.4 Numerical Experiments

First experiment: Classification

Write down bilevel classification problem and (if needed) which specification of the inexact bundle algorithm is used.

Write down the sets were used and how they were prepared.

References

- [1] P. Apkarian, D. Noll, and O. Prot. A trust region spectral bundle method for non-convex eigenvalue optimization. *SIAM Journal on Optimization*, 19(1):281–306, jan 2008.
- [2] Adil Bagirov, Napsu Karmitsa, and Marko M. Mäkelä. *Introduction to Nonsmooth Optimization: Theory, Practice and Software*. Springer International Publishing Switzerland, 2014.
- [3] Frank H. Clarke. *Optimization and nonsmooth analysis*. Classics in Applied Mathematics. Society for Industrial and Applied Mathematics Philadelphia, 1990.
- [4] A. Fuduli, M. Gaudioso, and G. Giallombardo. A dc piecewise affine model and a bundling technique in nonconvex nonsmooth minimization. *Optimization Methods and Software*, 19(1):89–102, 2004.
- [5] A. Fuduli, M. Gaudioso, and G. Giallombardo. Minimizing nonconvex nonsmooth functions via cutting planes and proximity control. *SIAM Journal on Optimization*, 14(3):743–756, 2004.
- [6] Napsu Haarala, Kaisa Miettinen, and Marko M. Mäkelä. Globally convergent limited memory bundle method for large-scale nonsmooth optimization. *Mathematical Programming*, 109(1):181–205, 2007.
- [7] Warren Hare and Claudia Sagastizábal. A redistributed proximal bundle method for nonconvex optimization. *SIAM Journal on Optimization*, 20(5):2442–2473, 2010.
- [8] Jean-Baptiste Hiriart-Urruty and Claude Lemaréchal. *Convex Analysis and Minimization Algorithms II*, volume 306 of *Grundlehren der mathematischen Wissenschaften*. Springer Berlin Heidelberg, 1993.
- [9] Jean-Baptiste Hiriart-Urruty and Claude Lemaréchal. *Convex Analysis and Minimization Algorithms I*, volume 305 of *Grundlehren der mathematischen Wissenschaften*. Springer Berlin Heidelberg, 2 edition, 1996.
- [10] Krzysztof C. Kiwiel. An aggregate subgradient method for nonsmooth and nonconvex minimization. *Journal of Computational and Applied Mathematics*, 14(3):391–400, 1986.
- [11] Claude Lemaréchal. Nonsmooth optimization and descent methods. Iiasa research report, International Institute for Applied Systems Analysis, 1978.
- [12] A. S. Lewis and S. J. Wright. A proximal method for composite minimization. *Mathematical Programming*, 158(1-2):501–546, aug 2015.
- [13] Marko M. Mäkelä and Pekka Neittaanmäki. *Nonsmooth Optimization: Analysis and Algorithms with Applications to Optimal Control*. World Scientific Pub Co Pte Lt, 1992.
- [14] Robert Mifflin. Semismooth and semiconvex functions in constrained optimization. *SIAM Journal on Control and Optimization*, 15(6):959–972, 1977.

- [15] Robert Mifflin. A modification and an extension of lemaréchal’s algorithm for nonsmooth minimization. In *Mathematical Programming Studies*, volume 17, pages 77–90. Springer Nature, 1982.
- [16] Robert Mifflin and Claudia Sagastizàbal. A science fiction story in nonsmooth optimization originating at iiasa. *Documenta Mathematica*, Extra Volume ISMP:291–300, 2012.
- [17] Yurii Nesterov and Vladimir Shikhman. Algorithmic principle of least revenue for finding market equilibria. In Boris Goldengorin, editor, *Optimization and Its Applications in Control and Data Sciences*, volume 115 of *Springer Optimization and Its Applications*, pages 381–435. Springer Nature, 2016.
- [18] Dominikus Noll. Cutting plane oracles to minimize non-smooth non-convex functions. *Set-Valued and Variational Analysis*, 18(3-4):531–568, sep 2010.
- [19] Dominikus Noll. Bundle method for non-convex minimization with inexact subgradients and function values. In *Computational and Analytical Mathematics*, pages 555–592. Springer Nature, 2013.
- [20] Dominikus Noll and Pierre Apkarian. Spectral bundle method for non-convex maximum eigenvalue functions: first-order methods. *Mathematical Programming*, 104(2-3):701–727, jul 2005.
- [21] Dominikus Noll, Olivier Prot, and Aude Rondepierre. A proximity control algorithm to minimize non-smooth and non-convex functions. *Pacific Journal of Optimization*, 4(3):571–604, 2012.
- [22] R. Tyrrell Rockafellar and Roger J. B. Wets. *Variational Analysis*, volume 317 of *Grundlehren der mathematischen Wissenschaften*. Springer Berlin Heidelberg, 3rd edition, 2009.
- [23] Helga Schramm and Jochem Zowe. A version of the bundle idea for minimizing a nonsmooth function: conceptual idea, convergence analysis, numerical results. *SIAM Journal on Optimization*, 2(1):121–152, feb 1992.
- [24] A. J. Smola, S.v.n. Vishwanathan, and V. Le Quoc. Bundle methods for machine learning. In J. C. Platt, D. Koller, Y. Singer, and S. T. Roweis, editors, *Advances in Neural Information Processing Systems*, number 20, 2007.
- [25] Choon Hui Teo, A. J. Smola, S.V.N. Vishwanathan, and Quoc V. Le. Bundle methods for regularized risk minimization. 2010.
- [26] Claudia Sagastizàbal Warren Hare. Computing proximal points of nonconvex functions. *Mathematical Programming*, 116:221–258, 2009.
- [27] Mikhail Solodov Warren Hare, Claudia Sagastizàbal. A proximal bundle method for nonsmooth nonconvex functions with inexact information. *Computational Optimization and Applications*, 63:1–28, 2016.